

计算机学院专业必修课

计算机组成

虚拟存储器

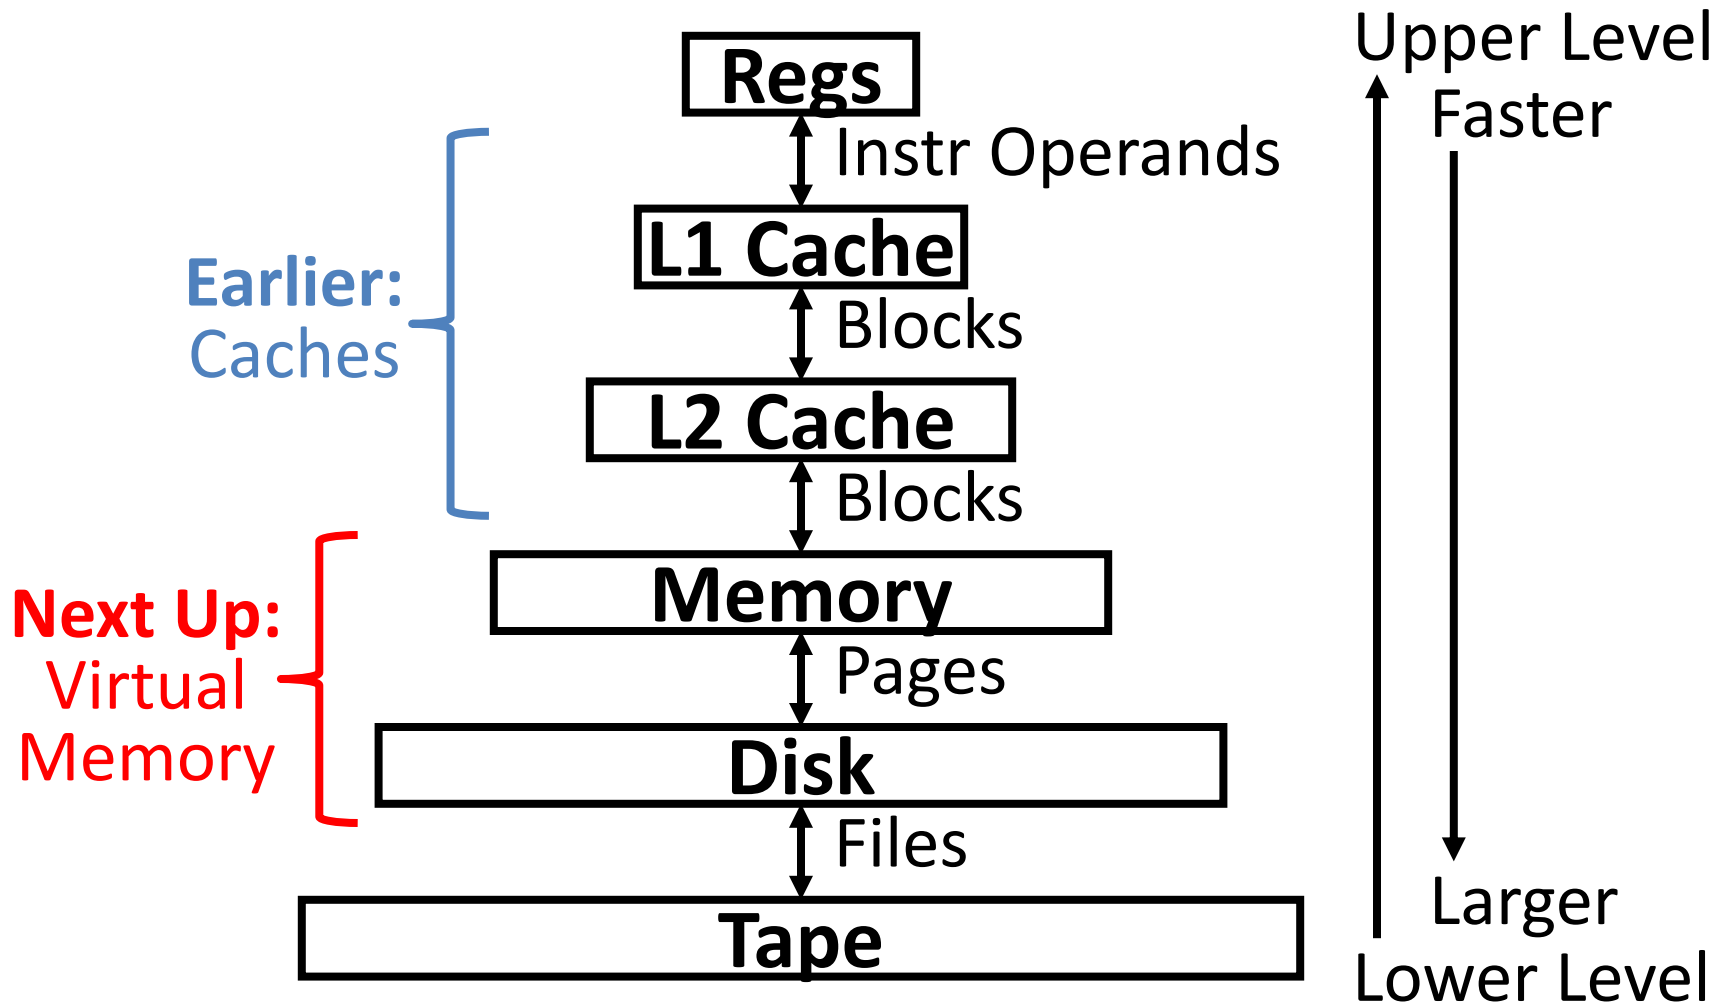
高小鹏

北京航空航天大学计算机学院

提纲

- 内容主要取材：CS61C的24讲和CS152的8讲
- 虚拟存储器(VM, Virtual Memory)
- 页表(Page Tables)
- TLB(Translation Lookaside Buffer)
- VM性能
- VM总结

Memory Hierarchy



Memory Hierarchy Requirements

- Principle of Locality
 - Allows caches to offer (close to) speed of cache memory with size of DRAM memory
 - Can we use this at the next level to give speed of DRAM memory with size of Disk memory?
- What other things do we need from our memory system?

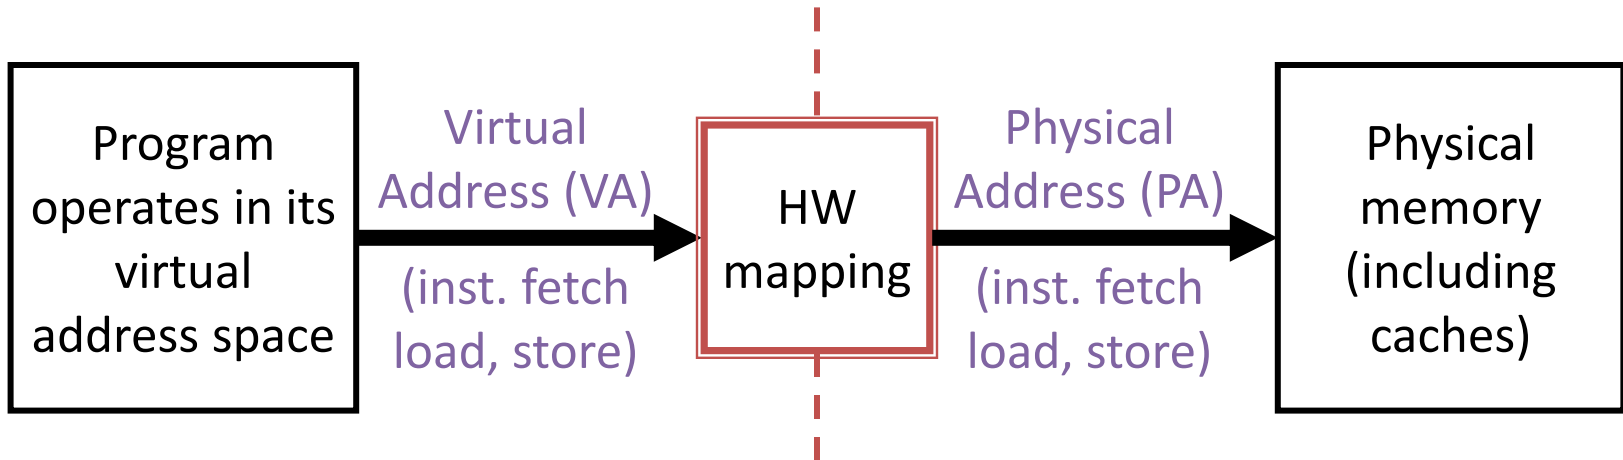
Memory Hierarchy Requirements

- Allow multiple processes to simultaneously occupy memory and provide *protection*
 - Don't let programs read from or write to each other's memories
- Give each program the illusion that it has its own *private address space* (via *translation*)
 - Suppose code starts at address 0x00400000, then different processes each think their code resides at the same address
 - Each program must have a different view of memory

Virtual Memory

- Next level in the memory hierarchy
 - Provides illusion of very large main memory
 - Working set of “pages” residing in main memory (subset of all pages residing on disk)
- **Main goal:** Avoid reaching all the way back to disk as much as possible
- **Additional goals:**
 - Let OS share memory among many programs and protect them from each other
 - Each process thinks it has all the memory to itself

Virtual to Physical Address Translation



- Each program operates in its own virtual address space and thinks it's the only program running
- Each is protected from the other
- OS can decide where each goes in memory
- Hardware gives virtual → physical mapping

VM Analogy (1/2)

- Trying to find a book in the UCB system
- Book title is like *virtual address (VA)*
 - What you want/are requesting
- Book call number is like *physical address (PA)*
 - Where it is actually located
- Card catalogue is like a *page table (PT)*
 - Maps from book title to call number
 - Does not contain the actual data you want
 - The catalogue itself takes up space in the library

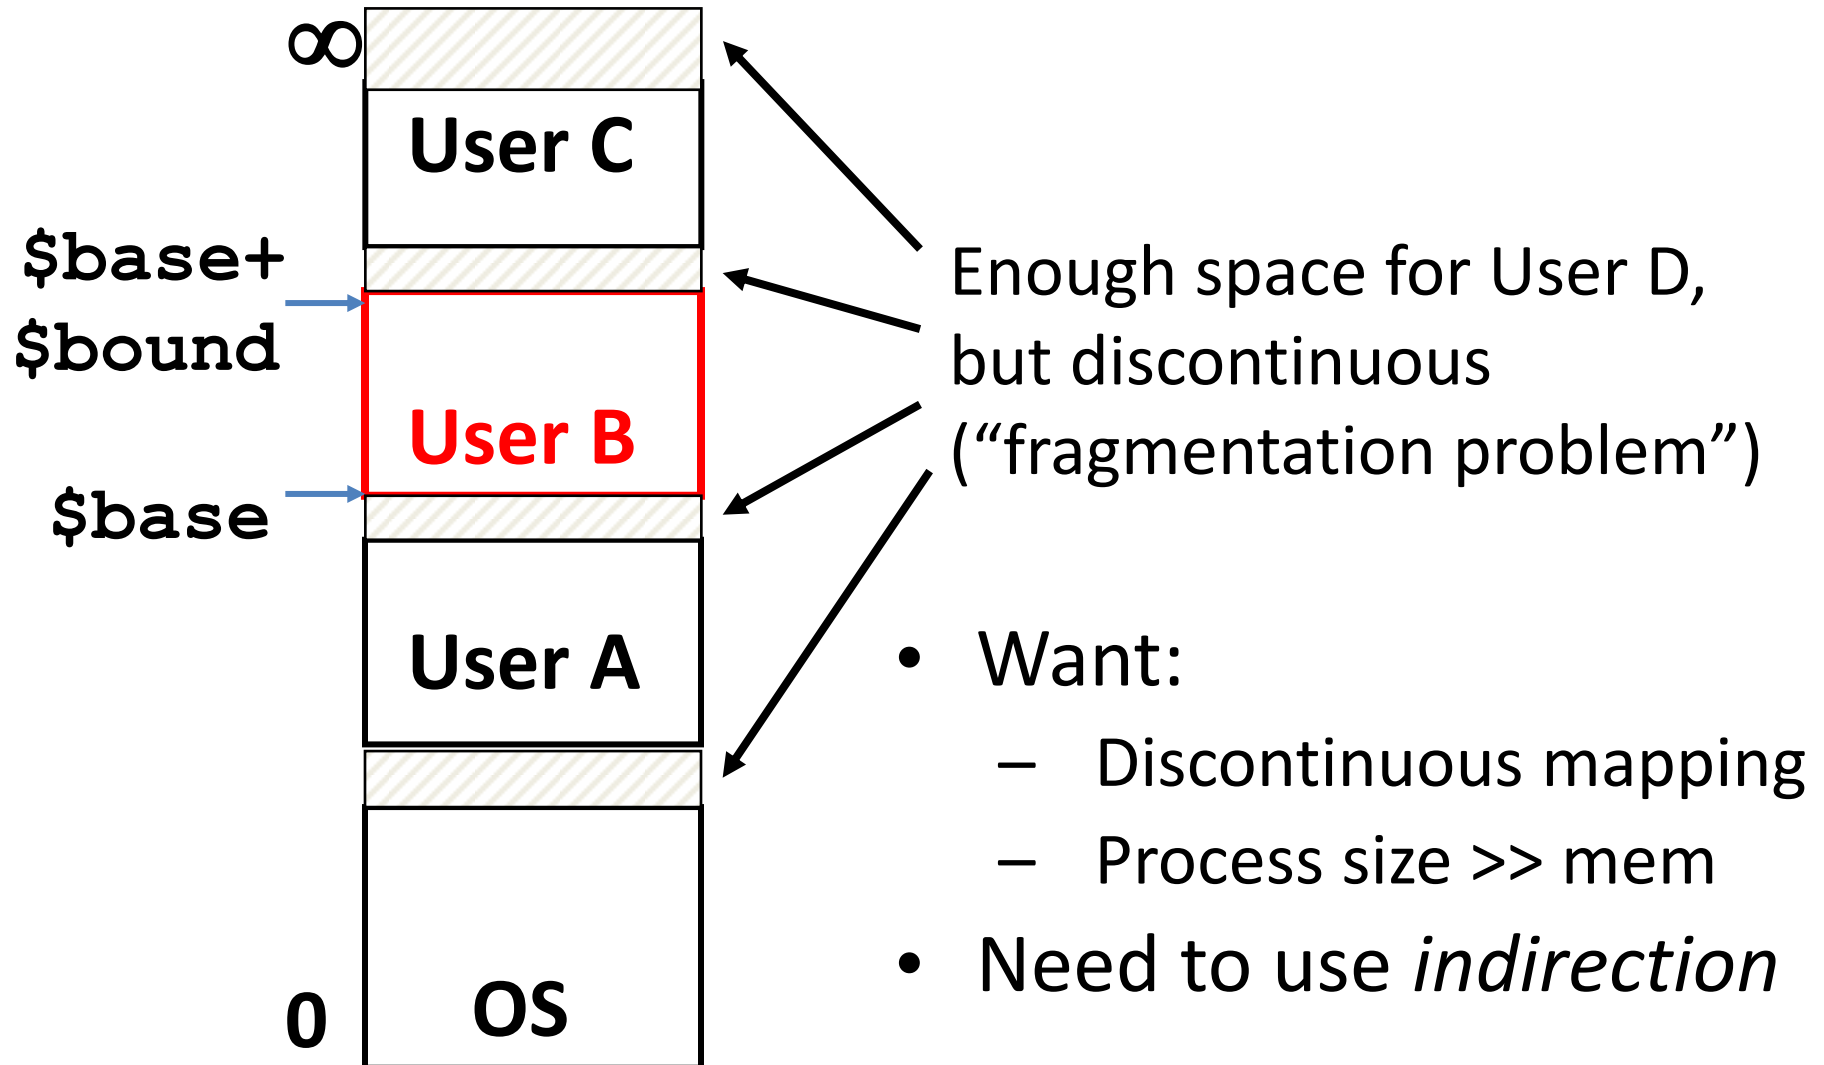
VM Analogy (2/2)

- Indication of current location within the library system is like *valid bit*
 - Valid if in current library (main memory) vs. invalid if in another branch (disk)
 - Found on the card in the card catalogue
- Availability/terms of use like *access rights*
 - What you are allowed to do with the book (ability to check out, duration, etc.)
 - Also found on the card in the card catalogue

提纲

- 内容主要取材：CS61C的24讲
- 虚拟存储器
- 页表(Page Tables)
- TLB(Translation Lookaside Buffer)
- VM性能
- VM总结

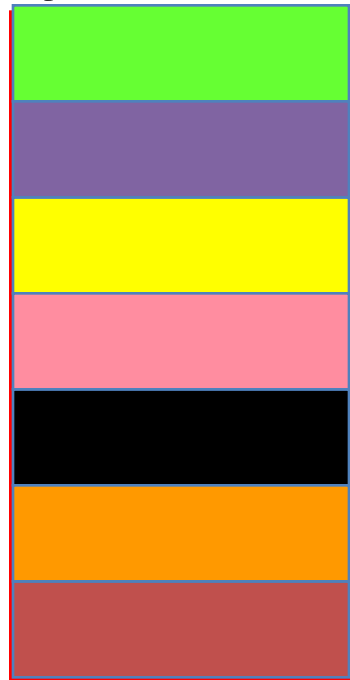
First Attempt: Base and Bound Reg



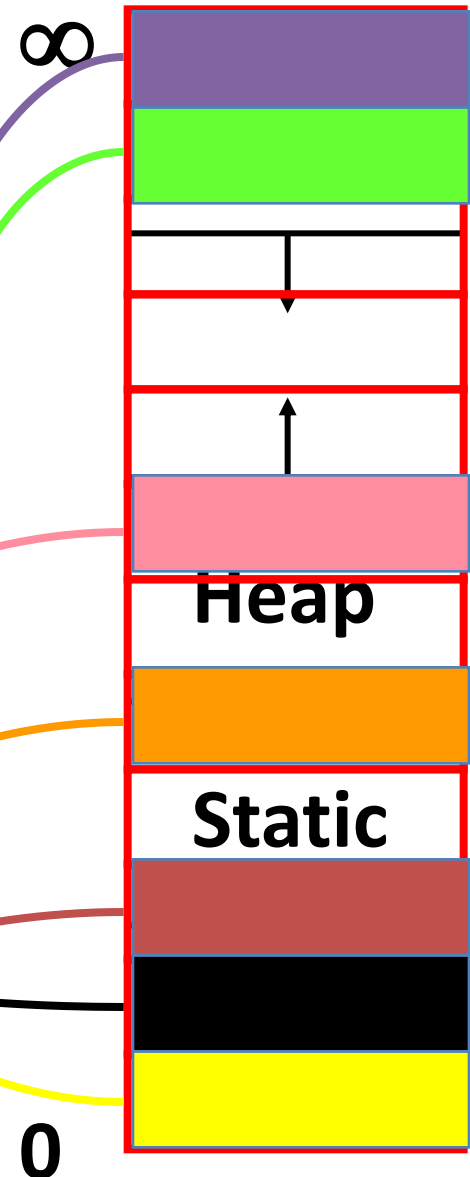
Mapping VM to PM

- Divide into equal sized chunks (about 4 KiB - 8 KiB)
- Any chunk of Virtual Memory can be assigned to any chunk of Physical Memory (“*page*”)

64 MB **Physical Memory**



Virtual Memory



Paging Organization

- Here assume page size is 4 KiB
 - Page is both unit of mapping and unit of transfer between disk and physical memory

Physical
Address

0x0000 page 0 4 Ki
0x1000 page 1 4 Ki
...
0x7000 page 7 4 Ki

Physical
Memory

Addr
Trans
MAP

Virtual
Address

0x0000 page 0 4 Ki
0x01000 page 1 4 Ki
0x02000 page 2 4 Ki
...
0x1F000 page 31 4 Ki

Virtual
Memory

Virtual Memory Mapping Function

- How large is main memory? Disk?
 - Don't know! Designed to be interchangeable components
 - Need a system that works regardless of sizes
- Use lookup table (*page table*) to deal with arbitrary mapping
 - Index lookup table by # of pages in VM (not all entries will be used/valid)
 - Size of PM will affect size of stored translation

Address Mapping

- Pages are aligned in memory
 - Border address of each page has same lowest bits
 - Page size is same in VM and PM, so denote lowest $O = \log_2(\text{page size}/B)$ bits as *page offset*
- Use remaining upper address bits in mapping
 - Tells you which page you want (similar to Tag)



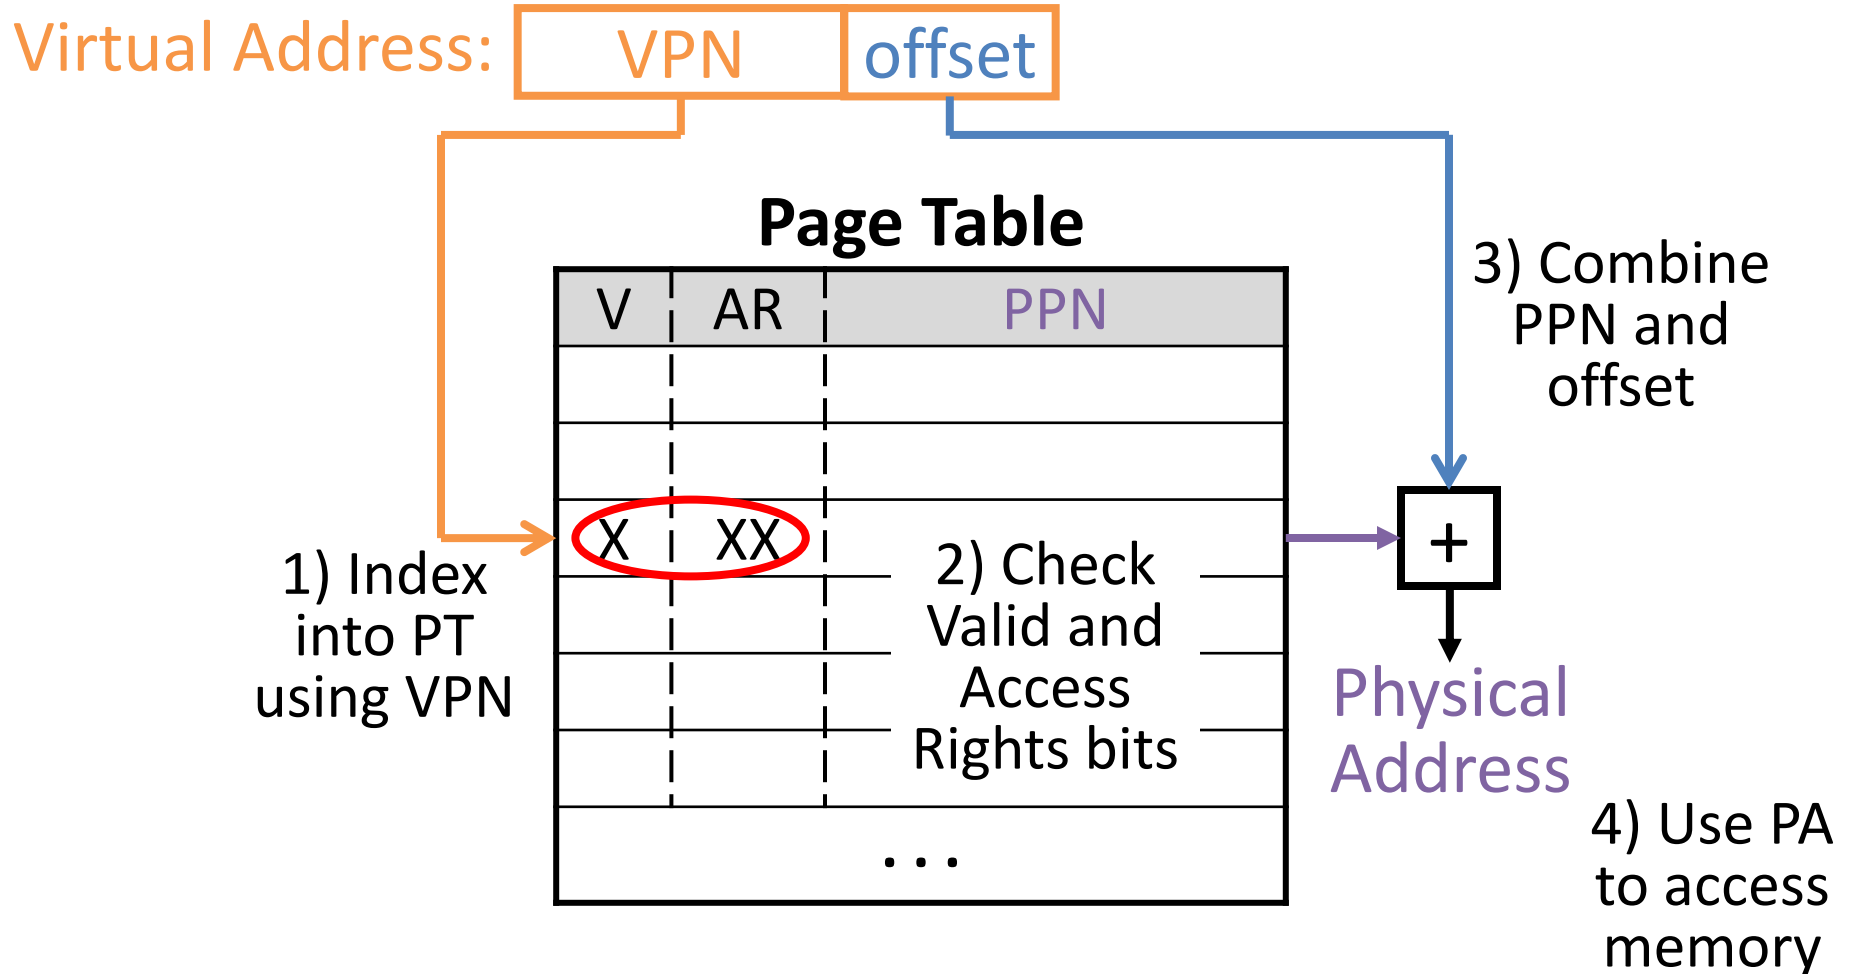
Address Mapping: Page Table

- **Page Table functionality:**
 - Incoming request is Virtual Address (**VA**), want Physical Address (**PA**)
 - Physical Offset = Virtual Offset (page-aligned)
 - So just swap Virtual Page Number (**VPN**) for Physical Page Number (**PPN**)

Physical Page # Virtual Page # / Page Offset

- **Implementation?**
 - Use VPN as index into PT
 - Store PPN and management bits (Valid, Access Rights)
 - Does NOT store actual data (the data sits in PM)

Page Table Layout



Page Table Entry Format

- Contains either PPN or indication not in main memory
- **Valid** = Valid page table entry
 - 1 → virtual page is in physical memory
 - 0 → OS needs to fetch page from disk
- **Access Rights** checked on every access to see if allowed (provides protection)
 - *Read Only*: Can read, but not write page
 - *Read/Write*: Read or write data on page
 - *Executable*: Can fetch instructions from page

Page Tables (1/2)

- A page table (PT) contains the mapping of virtual addresses to physical addresses
- Page tables located in main memory – Why?
 - Too large to fit in registers (2^{20} entries for 4 KiB pages)
 - Faster to access than disk and can be shared by multiple processors
- The OS maintains the PTs
 - Each process has its own page table
 - “State” of a process is PC, all registers, and PT
 - OS stores address of the PT of the *current* process in the *Page Table Base Register*

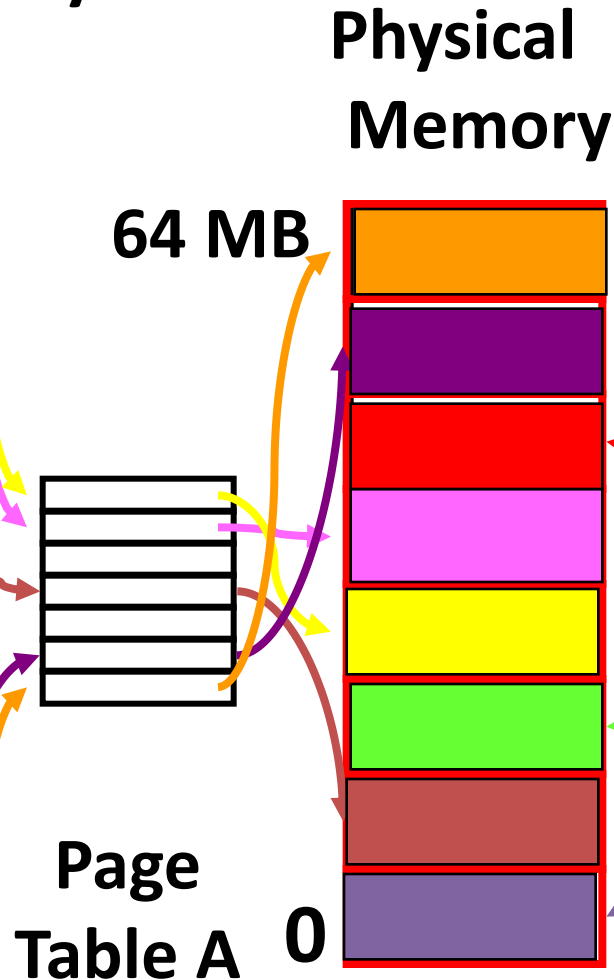
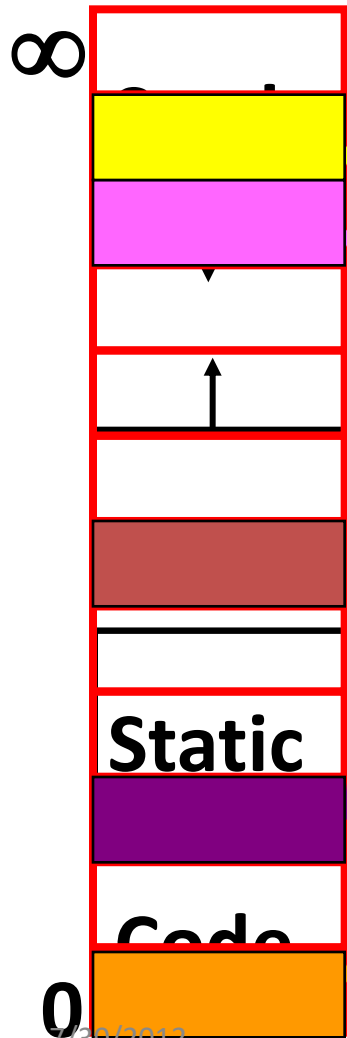
Page Tables (2/2)

- *Solves fragmentation problem*: all pages are the same size, so can utilize all available slots
- OS must reserve “*swap space*” on disk for *each* process
 - Running programs requires hard drive space!
- To grow a process, ask Operating System
 - If unused pages in PM, OS uses them first
 - If not, OS swaps some old pages (LRU) to disk

Paging/Virtual Memory Multiple Processes

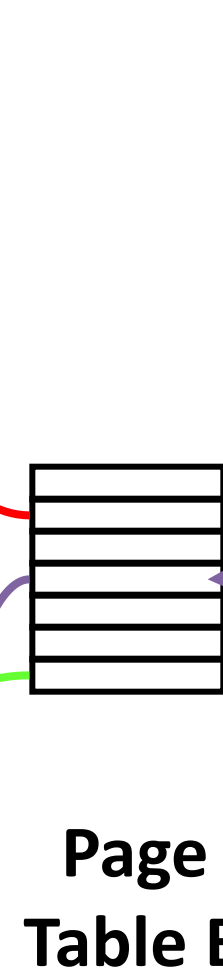
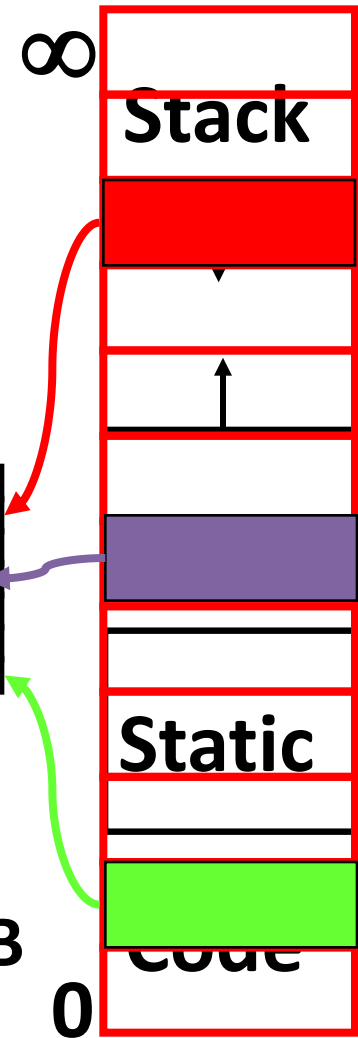
User A:

Virtual Memory



User B:

Virtual Memory



Review: Paging Terminology

- Programs use *virtual addresses (VAs)*
 - Space of all virtual addresses called *virtual memory (VM)*
 - Divided into pages indexed by *virtual page number (VPN)*
- Main memory indexed by *physical addresses (PAs)*
 - Space of all physical addresses called *physical memory (PM)*
 - Divided into pages indexed by *physical page number (PPN)*

Question: How many bits wide are the following fields?

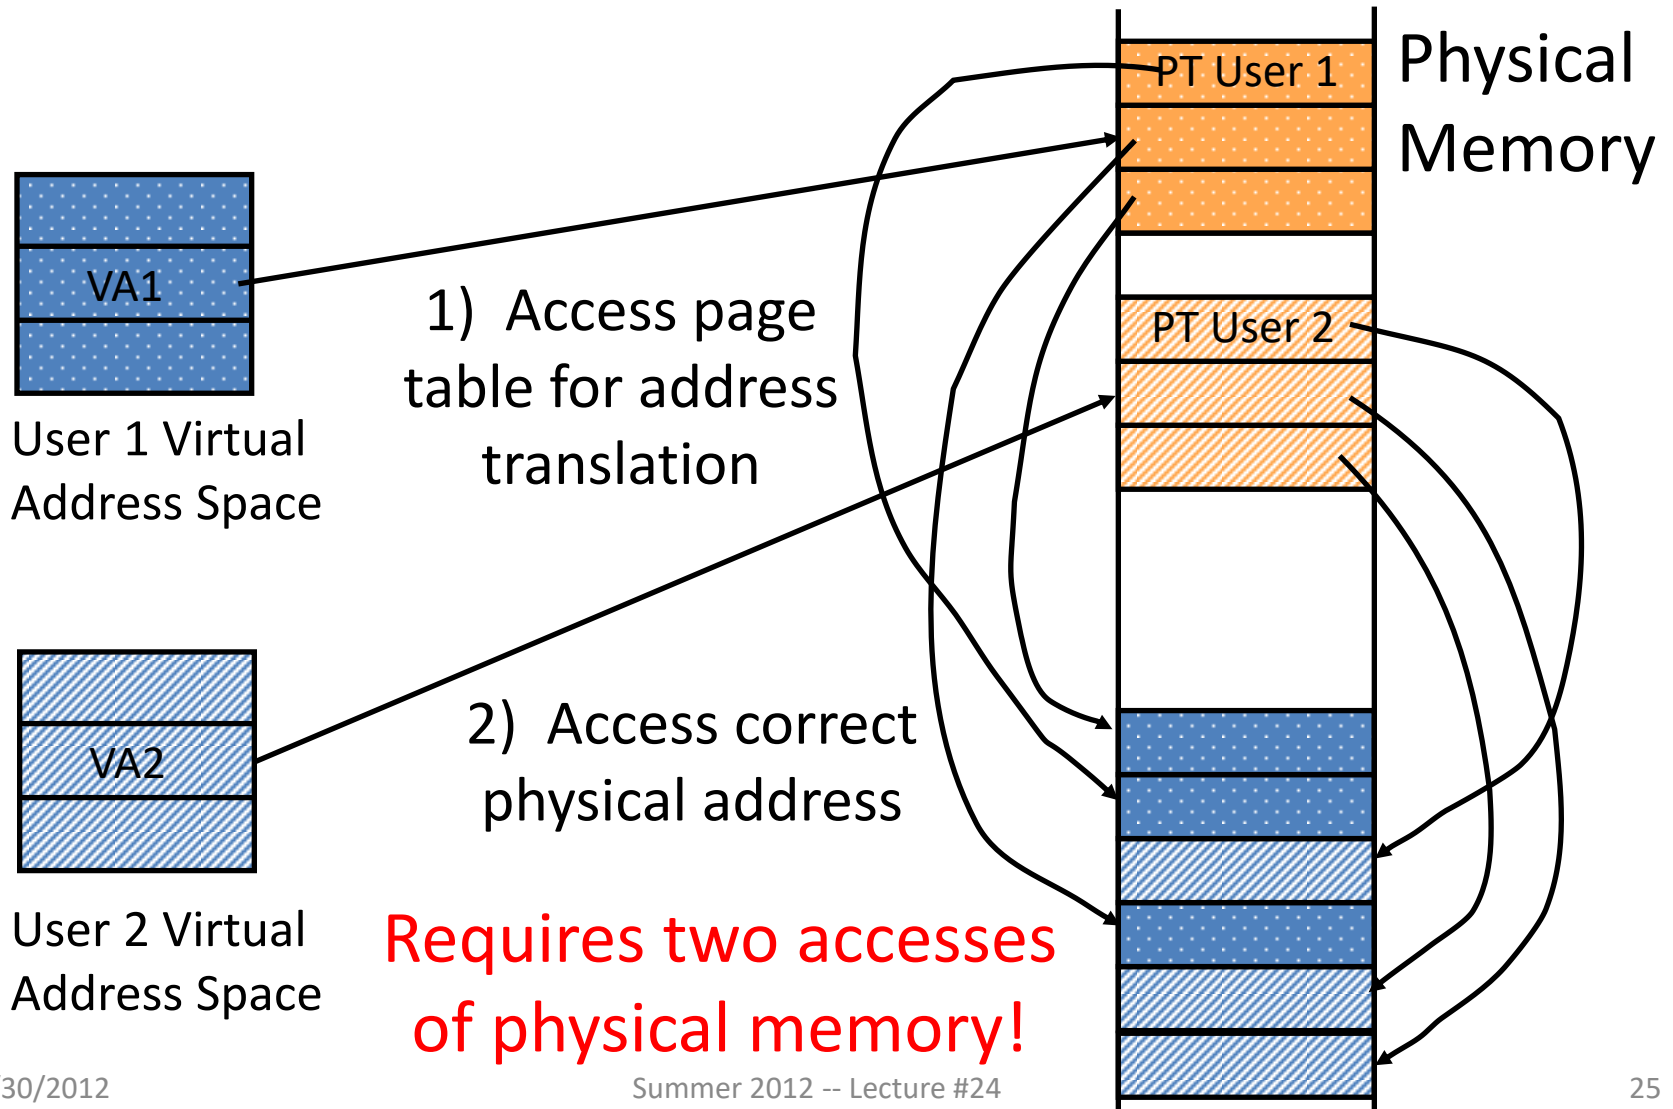
- 16 KiB pages
- 40-bit virtual addresses
- 64 GiB physical memory

	VPN	PPN
<input type="checkbox"/>	26	26
<input type="checkbox"/>	24	20
<input type="checkbox"/>	22	22
<input type="checkbox"/>	26	22

提纲

- 内容主要取材：CS61C的24讲
- 虚拟存储器
- 页表(Page Tables)
- TLB(Translation Lookaside Buffer)
- VM性能
- VM总结

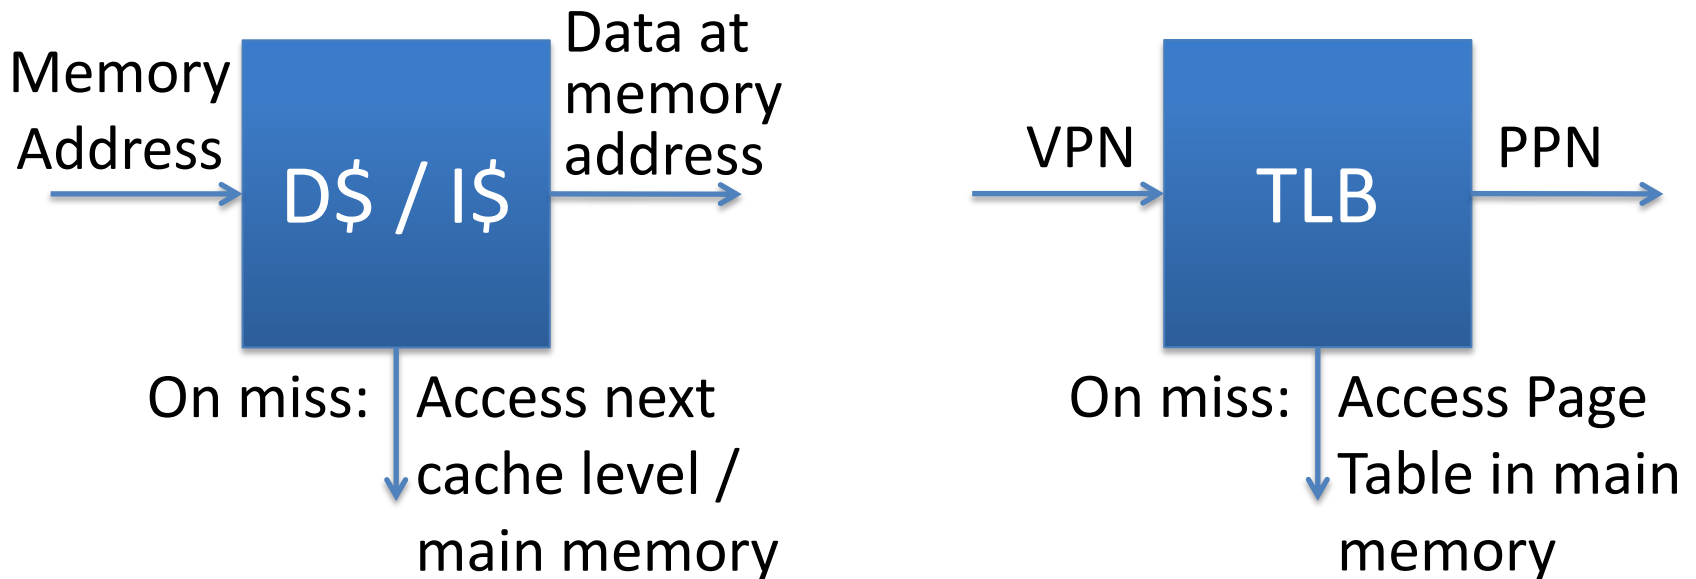
Retrieving Data from Memory



Virtual Memory Problem

- 2 physical memory accesses per data access
= SLOW!
- Since locality in pages of data, there must be locality in the translations of those pages
- Build a separate cache for the Page Table
 - For historical reasons, cache is called a *Translation Lookaside Buffer (TLB)*
 - Notice that what is stored in the TLB is NOT data, but the VPN \rightarrow PPN mapping translations

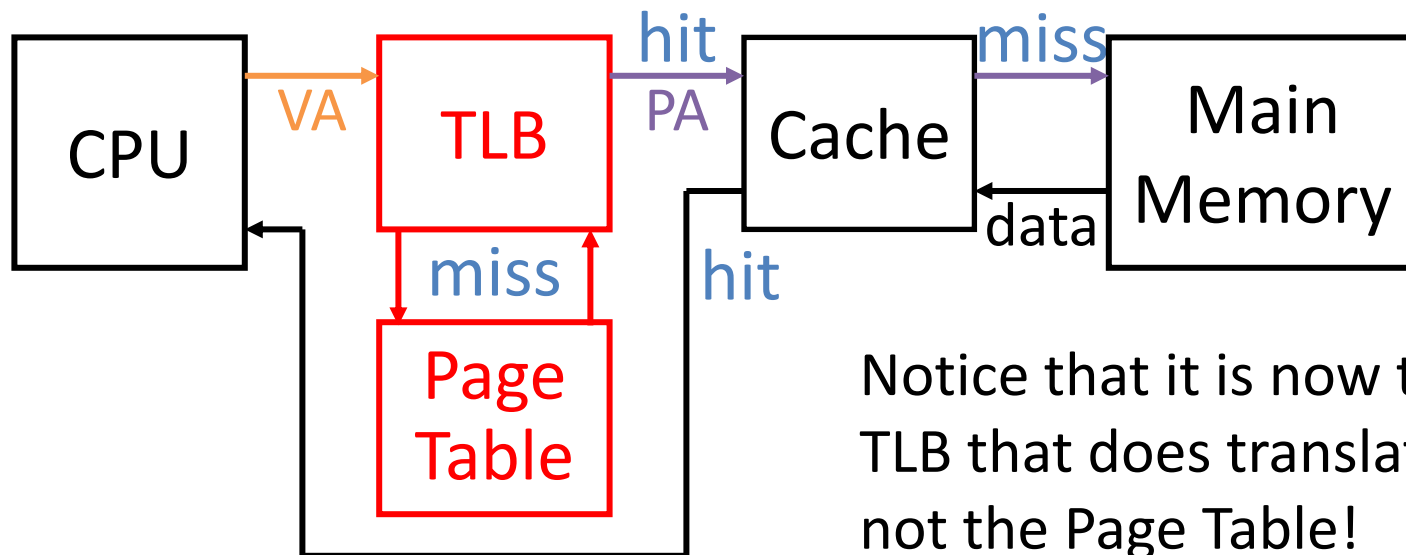
TLBs vs. Caches



- TLBs usually small, typically 16 – 512 entries
- TLB access time comparable to cache (« main memory)
- TLBs can have associativity
 - Usually fully/highly associative

Where Are TLBs Located?

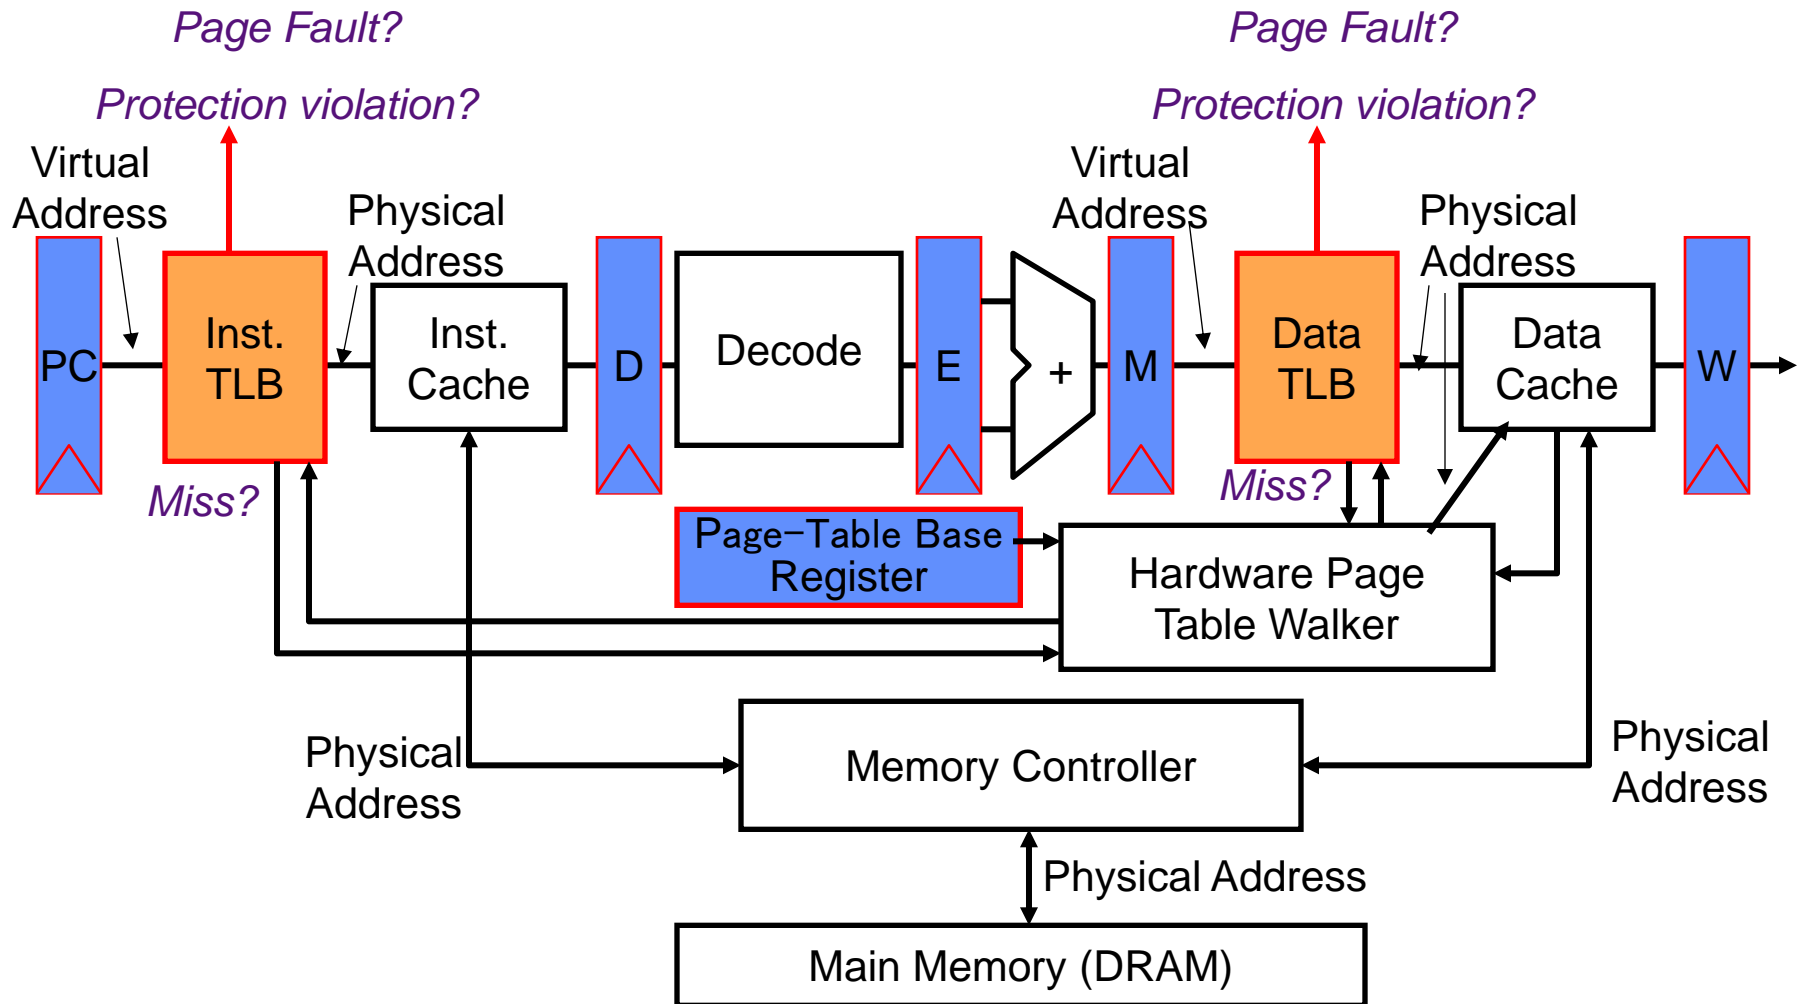
- Which should we check first: Cache or TLB?
 - Can cache hold requested data if corresponding page is not in physical memory? **No**
 - With TLB first, does cache receive VA or **PA**?



Notice that it is now the TLB that does translation, not the Page Table!

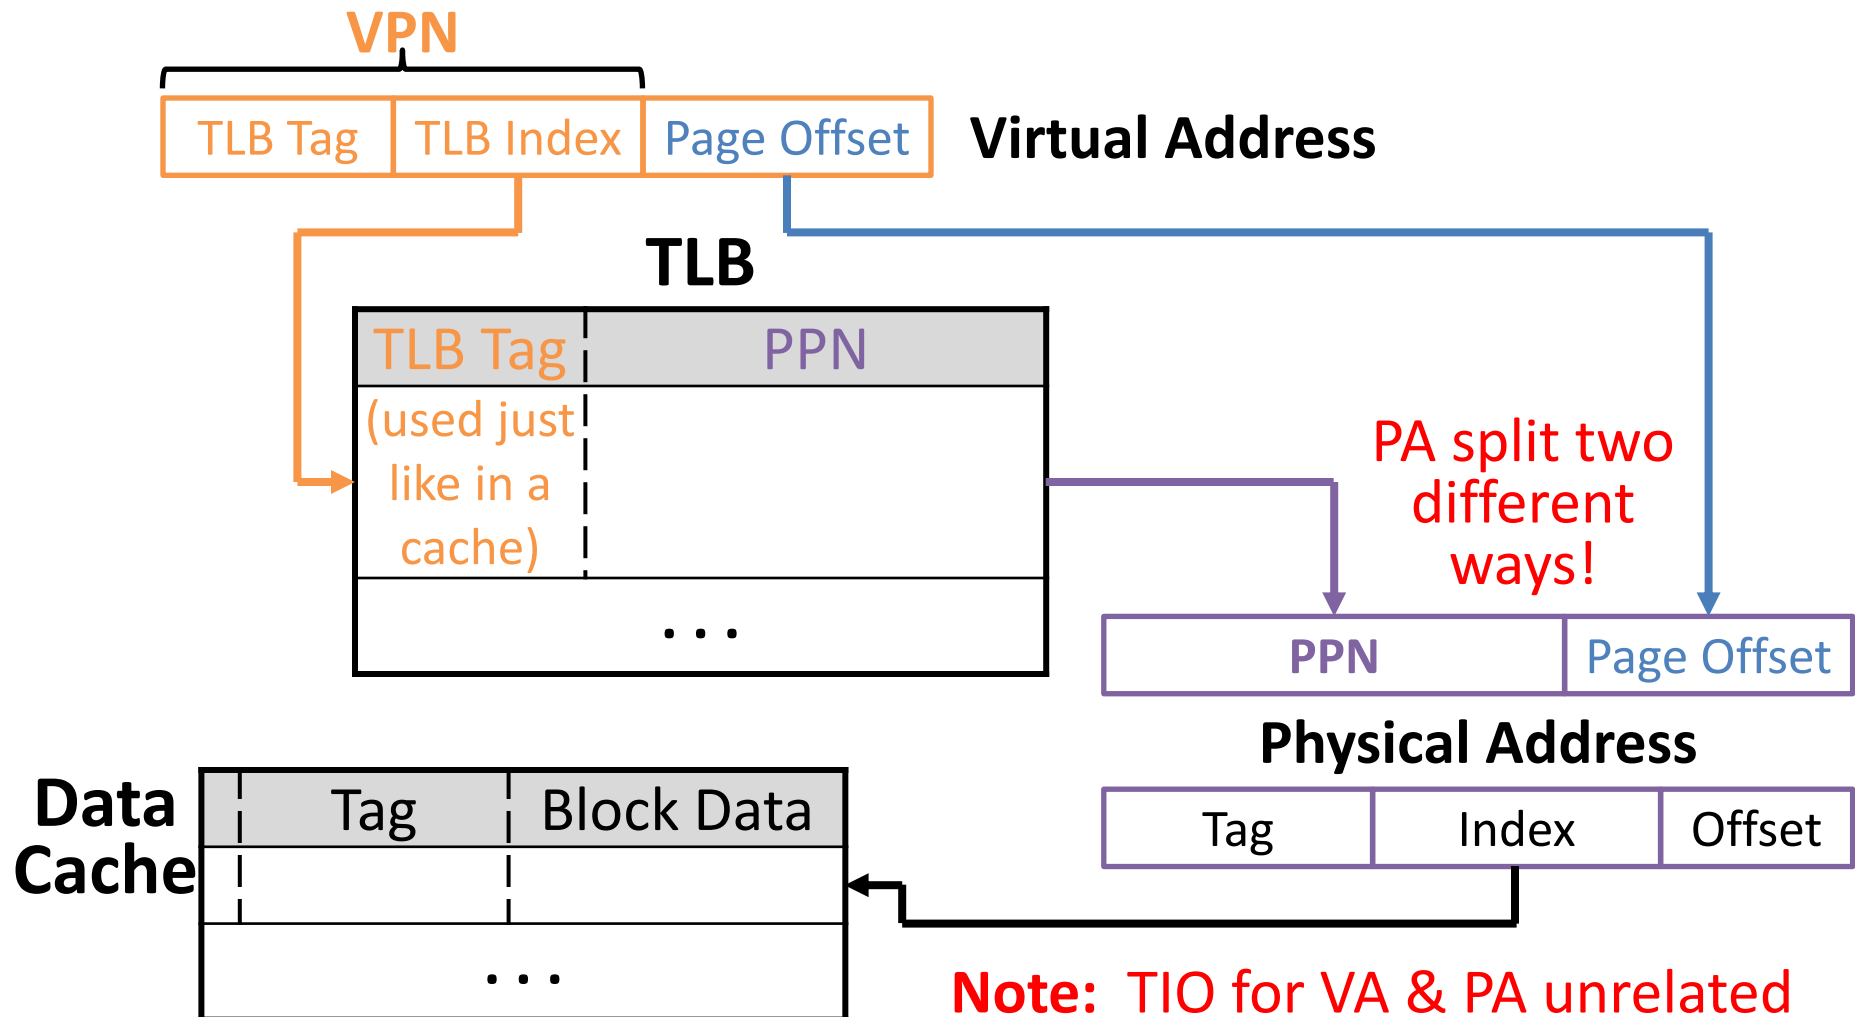


Page-Based Virtual-Memory Machine (Hardware Page-Table Walk)



- Assumes page tables held in untranslated physical memory

Address Translation Using TLB



Typical TLB Entry Format

Valid	Dirty	Ref	Access Rights	TLB Tag	PPN
X	X	X	XX		

- *Valid* and *Access Rights*: Same usage as previously discussed for page tables
- *Dirty*: Basically always use write-back, so indicates whether or not to write page to disk when replaced
- *Ref*: Used to implement LRU
 - Set when page is accessed, cleared periodically by OS
 - If Ref = 1, then page was referenced recently
- *TLB Tag*: VPN mod (# TLB entries)

Question: How many bits wide are the following?

- 16 KiB pages
- 40-bit virtual addresses
- 64 GiB physical memory
- 2-way set associative TLB with 512 entries

Valid	Dirty	Ref	Access Rights	TLB Tag	PPN
X	X	X	XX		

	TLB Tag	TLB Index	TLB Entry
<input type="checkbox"/>	12	14	38
<input type="checkbox"/>	18	8	45
<input type="checkbox"/>	14	12	40
<input type="checkbox"/>	17	9	43

Fetching Data on a Memory Read

- 1) Check TLB (input: VPN, output: PPN)
 - *TLB Hit*: Fetch translation, return PPN
 - *TLB Miss*: Check page table (in memory)
 - *Page Table Hit*: Load page table entry into TLB
 - *Page Table Miss* (*Page Fault*): Fetch page from disk to memory, update corresponding page table entry, then load entry into TLB
- 2) Check cache (input: PPN, output: data)
 - *Cache Hit*: Return data value to processor
 - *Cache Miss*: Fetch data value from memory, store it in cache, return it to processor

Page Faults

- Load the page off the disk into a free page of memory
 - Switch to some other process while we wait
- Interrupt thrown when page loaded and the process' page table is updated
 - When we switch back to the task, the desired data will be in memory
- If memory full, replace page (LRU), writing back if necessary, and update *both* page tables
 - Continuous swapping between disk and memory called “thrashing”

Performance Metrics

- VM performance also uses Hit/Miss Rates and Miss Penalties
 - *TLB Miss Rate*: Fraction of TLB accesses that result in a TLB Miss
 - *Page Table Miss Rate*: Fraction of PT accesses that result in a page fault
- Caching performance definitions remain the same
 - Somewhat independent, as TLB will always pass PA to cache regardless of TLB hit or miss

Data Fetch Scenarios

- Are the following scenarios for a single data access possible?
 - TLB Miss, Page Fault Yes
 - TLB Hit, Page Table Hit No
 - TLB Miss, Cache Hit Yes
 - Page Table Hit, Cache Miss Yes
 - Page Fault, Cache Hit No

Question: A program tries to load a word at X that causes a TLB miss but not a page fault. Are the following statements TRUE or FALSE?

- 1) The page table does not contain a valid mapping for the virtual page corresponding to the address X
- 2) The word that the program is trying to load is present in physical memory

	1	2
<input type="checkbox"/>	F	F
<input type="checkbox"/>	F	T
<input type="checkbox"/>	T	F
<input type="checkbox"/>	T	T

Updating Scenarios

- Using V = valid, D = dirty, R = ref to mean that field is set to the shown value for *any* entry in either PT or TLB
- Which of the following scenarios for a single data access are possible?
 - Read, D = 1 No
 - Write, R = 1 Yes
 - Read, V = 0 Yes
 - Write, D = 0 No

Question: Assume the page table entry in question is present in the TLB and we are using a uniprocessor system. Are the following statements TRUE or FALSE?

- 1) The valid bit for that page must be the same in the PT and TLB
- 2) The dirty bit for that page must be the same in the PT and TLB

	1	2
<input type="checkbox"/>	F	F
<input type="checkbox"/>	F	T
<input type="checkbox"/>	T	F
<input type="checkbox"/>	T	T

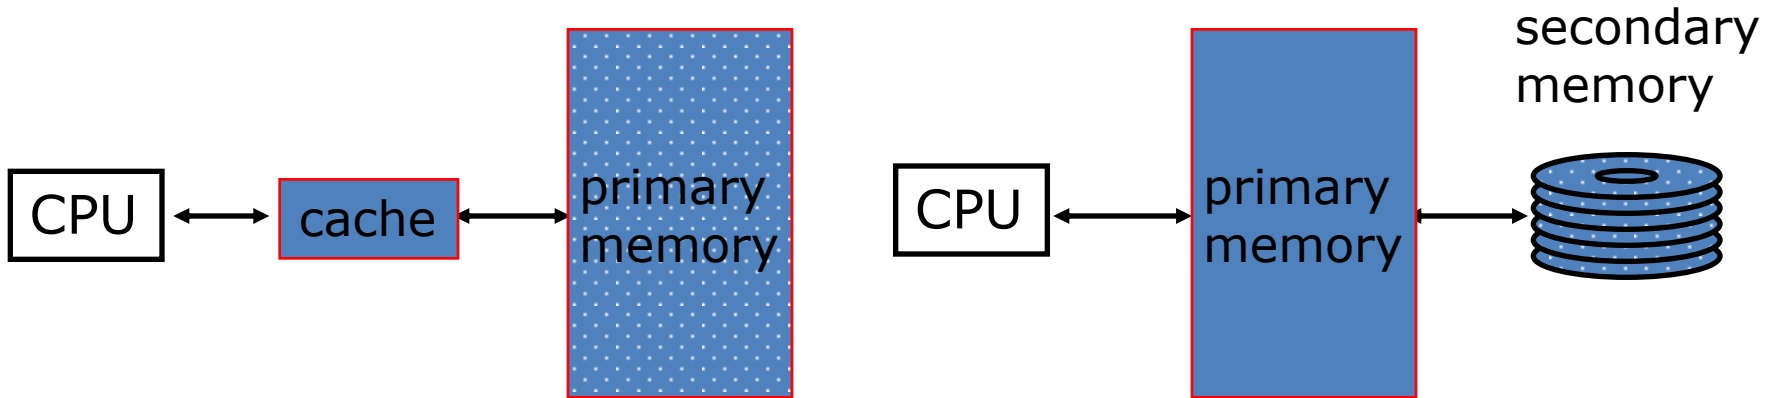
提纲

- 内容主要取材：CS61C的24讲
- 虚拟存储器
- 页表(Page Tables)
- TLB(Translation Lookaside Buffer)
- VM性能
- VM总结

VM Performance

- Virtual Memory is the level of the memory hierarchy that sits *below* main memory
 - TLB comes *before* cache, but affects transfer of data from disk to main memory
 - Previously we assumed main memory was lowest level, now we just have to account for disk accesses
- Same CPI, AMAT equations apply, but now treat main memory like a mid-level cache

Typical Performance Stats



Caching

- cache entry
- cache block (≈ 32 bytes)
- cache miss rate (1% to 20%)
- cache hit (≈ 1 cycle)
- cache miss (≈ 100 cycles)

Demand paging

- page frame
- page (≈ 4 Ki bytes)
- page miss rate ($< 0.001\%$)
- page hit (≈ 100 cycles)
- page miss (≈ 5 M cycles)

Impact of Paging on AMAT (1/2)

- Memory Parameters:
 - L1 cache hit = 1 clock cycles, hit 95% of accesses
 - L2 cache hit = 10 clock cycles, hit 60% of L1 misses
 - DRAM = 200 clock cycles (≈ 100 nanoseconds)
 - Disk = 20,000,000 clock cycles (≈ 10 milliseconds)
- Average Memory Access Time (no paging):
 - $1 + 5\% \times 10 + 5\% \times 40\% \times 200 = 5.5$ clock cycles
- Average Memory Access Time (with paging):
 - 5.5 (AMAT with no paging) + ?

Impact of Paging on AMAT (2/2)

- Average Memory Access Time (with paging) =
 - $5.5 + 5\% \times 40\% \times (1 - \text{HR}_{\text{Mem}}) \times 20,000,000$
- AMAT if $\text{HR}_{\text{Mem}} = 99\%$?
 - $5.5 + 0.02 \times 0.01 \times 20,000,000 = 4005.5$ ($\approx 728x$ slower)
 - 1 in 20,000 memory accesses goes to disk: 10 sec program takes 2 hours!
- AMAT if $\text{HR}_{\text{Mem}} = 99.9\%$?
 - $5.5 + 0.02 \times 0.001 \times 20,000,000 = 405.5$
- AMAT if $\text{HR}_{\text{Mem}} = 99.9999\%$
 - $5.5 + 0.02 \times 0.000001 \times 20,000,000 = 5.9$

提纲

- 内容主要取材：CS61C的24讲
- 虚拟存储器
- 页表(Page Tables)
- TLB(Translation Lookaside Buffer)
- VM性能
- VM总结

Virtual Memory Motivation

- Memory as cache for disk (reduce disk accesses)
 - Disk is so slow it significantly affects performance
 - Paging maximizes memory usage with large, evenly-sized pages that can go anywhere
- Allows processor to run multiple processes simultaneously
 - Gives each process illusion of its own (large) VM
 - Each process uses standard set of VAs
 - Access rights provide *protection*

Paging Summary

- Paging requires address *translation*
 - Can run programs larger than main memory
 - Hides variable machine configurations (RAM/HD)
 - Solves fragmentation problem
- Address mappings stored in page tables in memory
 - Additional memory access mitigated with TLB
 - Check TLB, then Page Table (if necessary), then Cache

Hardware/Software Support for Memory Protection

- Different tasks can share parts of their virtual address spaces
 - But need to protect against errant access
 - Requires OS assistance
- Hardware support for OS protection
 - Privileged supervisor mode (a.k.a. *kernel mode*)
 - Privileged instructions
 - Page tables and other state information only accessible in supervisor mode
 - System call exception (e.g. `syscall` in MIPS)

Context Switching

- How does a single processor run many programs at once?
- *Context switch*: Changing of internal state of processor (switching between processes)
 - Save register values (and PC) and change value in Page Table Base register
- What happens to the TLB?
 - Current entries are for different process
 - Set all entries to invalid on context switch

Virtual Memory Summary

- User program view:
 - Contiguous memory
 - Start from some set VA
 - “Infinitely” large
 - Is the only running program
- Reality:
 - Non-contiguous memory
 - Start wherever available memory is
 - Finite size
 - Many programs running simultaneously
- Virtual memory provides:
 - Illusion of contiguous memory
 - All programs starting at same set address
 - Illusion of \sim infinite memory (2^{32} or 2^{64} bytes)
 - Protection, Sharing
- Implementation:
 - Divide memory into chunks (pages)
 - OS controls page table that maps virtual into physical addresses
 - memory as a cache for disk
 - TLB is a cache for the page table

